

ECE 661 Homework 7

Suyash Ail
sail@purdue.edu

October 20, 2020

The task of this homework is to implement a simple image classifier using the Local Binary Pattern (LBP) feature extraction and Nearest Nearest (NN) classifier.

1 Theory Questions

Question1

This is a very good chart explaining all the texture based characterization techniques:

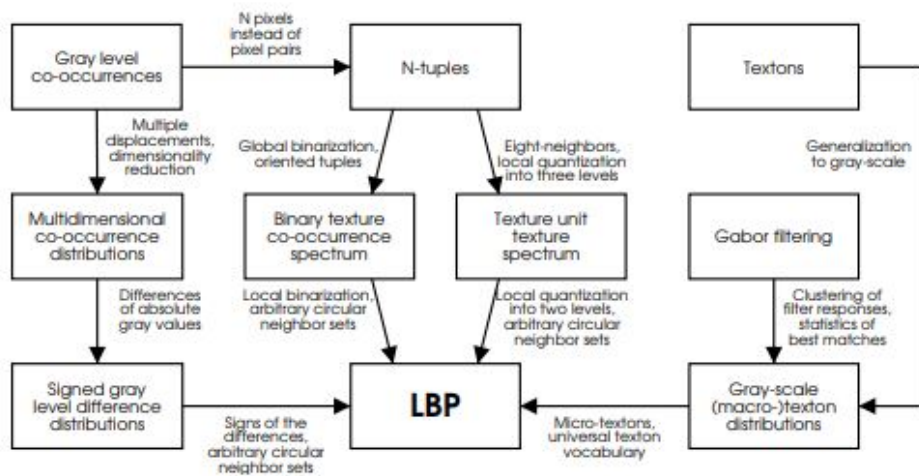


Fig. 2. LBP in the field of texture analysis operators.

Figure 1: Texture based techniques [1]

1.0.1 Gray Level Co-Occurrence Matrix (GLCM)

- In case of (GLCM), we try to estimate the joint probability distribution $P[x_1, x_2]$ for grayscale values in an image where x_1 is grayscale value of any randomly selected pixel in the image and x_2 is the grayscale value of another pixel at a specific vector distance d from x_1
- We examine grayscale values at two different pixels that are separated by a distance vector d . After this, we count the number of pairs that are d distance apart but

exhibit the grayscale values (x_1, x_2) . We normalize this count to get $P[x_1, x_2]$, we can characterize the texture based on the shape of $P[x_1, x_2]$

- Using this joint probability distribution (microscopic view of texture) we can compute values of Entropy, Energy, Contrast and Homogeneity (macroscopic view of textures) that help with texture characterizations

1.0.2 Local Binary Pattern (LBP)

LBP's compute a local representation of texture. This local representation is constructed by comparing each pixel with its surrounding neighborhood of pixels.

1. The fundamental idea of LBP is the notion of a binary pattern to characterize the grayscale variations around the pixels using 0s and 1s.
2. The pixels are assumed to lie on the circumference of a circle of radius R with the centre pixel as the center of the circle.
3. Not all points fall at pixel locations and hence bilinear interpolation is used to generate their pixel values.
4. Once these pixel values are obtained, we threshold them with the center pixel. If the center pixel is higher than the neighbour pixels, they are set to 0, else 1. Hence we obtain a P dimensional vector representation of the local texture of the image. The vector can be converted into an integer value which replaces the center pixel.
5. The binary pattern should be invariant of scale and rotation. This is obtained by circular shifting the feature vectors until the minimum integer representation is obtained.

The binary pattern and the texture represented by them can be visualized as below(White points represent 1s and black points represent 0s) :




Fig. 3. Different texture primitives detected by the LBP.

Figure 2: LBP patterns and corresponding textures [1]

1.0.3 Gabor Filter Family

Gabor filters are used to extract localized features using the convolution operators. It is a part of the MPEG-7 standard. The convolution kernels are Gaussian weighted 2D Fourier kernels.

$$h(x, y) = \frac{1}{2\pi\sigma_x\sigma_y} \exp\left(\frac{-1}{2} \left[\frac{x^2}{\sigma_x^2} + \frac{y^2}{\sigma_y^2}\right]\right) \cdot \cos 2\pi u_0 x$$

The Gaussian kernel helps in achieving localization while the Fourier transform helps in extracting spatial frequencies from the image. Multiple frequencies at multiple angles are used to extract an array of Gabor filtered channels which form the feature vector.

Question2

1. RGB and HSI are just linear variants of each other : True
2. The color space L*a*b* is a nonlinear model of color perception.: True
3. Measuring the true color of the surface of an object is made difficult by the spectral composition of the illumination.: True

2 Logic

2.1 LBP feature extraction

The LBP feature extraction is a texture based feature extraction technique which is used to compute translation and rotation invariant features.

2.1.1 Creating binary patterns for inter-pixel variations

First step is define the neighbourhood of each pixel X. The neighbourhood is defined by two parameters - radius (R) and number of neighbours(P).

The P neighbours lie along a circle of radius R. Their position can be computed by

$$(\Delta u, \Delta v) = (R \cos(\frac{2\pi p}{P}), R \sin(\frac{2\pi p}{P}))$$

where $p = 0, 1, 2, \dots, P$ and Δu and Δv indicate the displacement of the point p in the X and Y direction. Let A be the center pixel with coordinates (x,y). Then the neighbouring points are given by $x + \Delta u$ and $y + \Delta v$. $p = 0$ corresponds to the bottom pixel and $p = 4$ corresponds to the top pixel. All the P points will not lie on pixel coordinates. Hence to determine their location, we need bilinear interpolation.

$$I_p = (1 - \Delta u)(1 - \Delta v)A + (1 - \Delta u)(\Delta v)B + (\Delta u)(1 - \Delta v)C + (\Delta u)(\Delta v)D$$

where A,B,C,D represent the four corners of the rectangle around the center pixel.

Next, these intensity values are threshold with respect to the center pixel in order to obtain the binary pattern. If $I_p > A$, the value is 1 else 0.

2.1.2 Rotational Invariance in LBP

Given the binary pattern, the goal is to create rotation invariance patterns. To achieve this, the binary sequence is circularly shifted to obtain minimum integer when the binary sequence is converted to integer.

I have made use of professor Avinash Kak's BitVector module to obtain the binary sequence corresponding to minimum integer values.

2.1.3 Image Encoding

Next, the rotationally invariant binary pattern is encoded by a single integer ranging from 0 to $P+2$. The order of the binary pattern is used for encoding. The methodology is as follows:

1. If there are exactly two runs, a run of 0s followed by a run of 1s, then represent the pattern by number of ones in the sequence.
2. If the pattern has all 0s, represent the pattern by 0
3. If the pattern has all 1s, encode the pattern by 1
4. if the pattern has more than two runs, encode the pattern by $P+1$.

For each point, we compute the binary encoding using the above steps. One thing to make sure is that all images are of the same size. Or else the histogram for a high resolution image and a low resolution image of the same class would differ heavily. We wish to keep the method invariant of the size of the image. Hence the image is first resized to (256,256) and then the computed histogram is normalized.

2.2 K-Nearest Neighbours

KNN is a supervised classification algorithm that requires both the feature and their corresponding class labels. KNN requires a distance metric to evaluate the closeness between the two feature vectors. In this homework I have used Euclidean distance and Manhattan distance.

There is no training as such for the KNN algorithm. During the testing phase, each new instance is compared with every feature in the train dataset. The class labels corresponding to the K closest points are chosen to determine the final class using a voting method. The class that occurs maximum times is selected as the predicted class.

3 Results



Figure 3: Beach




Figure 4: Beach histogram



Figure 5: Building




Figure 6: Building histogram



Figure 7: car




Figure 8: Car histogram



Figure 9: Mountain




Figure 10: Mountain histogram



Figure 11: Tree




Figure 12: Tree histogram

3.1 Confusion Matrix

$$\text{ConfMat} = \begin{bmatrix} \text{True/Predicted} & \text{Tree} & \text{Beach} & \text{Mountain} & \text{Building} & \text{Car} \\ \text{Tree} & 3 & 0 & 1 & 1 & 0 \\ \text{Beach} & 0 & 5 & 0 & 0 & 0 \\ \text{Mountain} & 0 & 2 & 2 & 1 & 0 \\ \text{Building} & 1 & 0 & 1 & 2 & 1 \\ \text{Car} & 0 & 0 & 0 & 2 & 3 \end{bmatrix}$$

Accuracy of the classifier =

$$\frac{3 + 5 + 2 + 2 + 3}{25} = 0.6$$

The classifier has correctly classified all the Beach images. It performed well on tree and car but had trouble with mountain and building images.

Varying the distance metric did not change the accuracy. By increasing the K to 7, i got a higher accuracy.

$$\text{ConfMat} = \begin{bmatrix} \text{True/Predicted} & \text{Tree} & \text{Beach} & \text{Mountain} & \text{Building} & \text{Car} \\ \text{Tree} & 3 & 0 & 1 & 1 & 0 \\ \text{Beach} & 0 & 5 & 0 & 0 & 0 \\ \text{Mountain} & 0 & 1 & 3 & 1 & 0 \\ \text{Building} & 1 & 0 & 1 & 2 & 1 \\ \text{Car} & 0 & 0 & 0 & 2 & 3 \end{bmatrix}$$

Accuracy of the classifier =

$$\frac{3 + 5 + 3 + 3 + 3}{25} = 0.68$$

Further increasing the K value resulted in decrease of accuracy.

4 References

[1] Maenpaa, Topi Pietikainen, Matti. (2005). Texture analysis with local binary patterns. Handbook of Pattern Recognition and Computer Vision. 10.1142/9789812775320_011.

5 Code Listings

```
1 # -*- coding: utf-8 -*-
2 """hw7_suyash_ail.ipynb
3
4 Automatically generated by Colaboratory.
5
6 Original file is located at
7     https://colab.research.google.com/drive/1uRQot1tPBL6wfsXMYHkGMDVlxXuJVxSB
8 """
9
10 import os
11 os.chdir( '/content/drive/My Drive/BitVector-3.4.9/BitVector ' )
12
13 import numpy as np
14 import matplotlib.pyplot as plt
15 import glob
16 import cv2
17 from BitVector import BitVector
18 import pickle
19
20 '''
21 def n_neighbour(image, radius=1, num_neighbours=8):
22     R = radius
23     P = num_neighbours
24
25     dx_list = []
26     dy_list = []
27     for p in range(P):
28         du = R*np.cos(2*np.pi*p/P) #get x-coord
29         dv = R*np.sin(2*np.pi*p/P) #get y-coord
30
31         if abs(du)< 1e-4: du=0
32         if abs(dv)< 1e-4: dv=0
33
34         dx_list.append(du)
35         dy_list.append(dv)
36
37     return dx_list, dy_list
38 '''
39
40 def bilinear_intp(A,B,C,D,dx,dy):
41     '''
42     Function that returns the bilinear interpolation of a point with respect to
43     its 4 neighbour points.
44
45     A-----B
46     |       |
47     |   X   |
48     |       |
49     C-----D
50
51     A is the CENTER pixel
52     dx : x distance from A
53     dy : y distance from A
54     '''
55     return (1-dy)*(1-dx)*A + (1-dy)*dx*B + dy*(1-dx)*C + dy*dx*D
56
57 def lbp (A):
58     # A is the n neighbourhood of the pixel
```

```

56     ' , '
57     a00  a01  a02
58     a10  'a11' a12
59     a20  a21  a22
60
61     V vector starts from from a21 and moves counter-clockwise
62     v6 v5 v4
63     v7     v3
64     v8 v1 v2
65
66     v1,v3,v5,v7 are the same as original pixels
67
68     for now hard-coded the distances..
69     ' , '
70     v1 = A[2,1]
71     v2 = bilinear_intp(A[1,1], A[1,2], A[2,1], A[2,2], 0.707, 0.707)
72     v3 = A[1,2]
73     v4 = bilinear_intp(A[1,1],A[1,2],A[0,1],A[0,2],0.707,0.707)
74     v5 = A[0,1]
75     v6 = bilinear_intp(A[1,1],A[1,0],A[0,1],A[0,0],0.707,0.707)
76     v7 = A[1,0]
77     v8 = bilinear_intp(A[1,1],A[1,0],A[2,1],A[2,0],0.707,0.707)
78
79     vec = np.array([v1,v2,v3,v4,v5,v6,v7,v8])
80     vec = vec >= A[1,1]
81
82
83     ##### Using prof. Avinash Kak's BitVector module to get the lbp value
84     #####
85     # https://pypi.org/project/BitVector/
86     bv = BitVector(bitlist = vec)
87     intvals = [int(bv<<1) for _ in range(P)] # integer values for circular
88     shift
89     #print('intvals=',intvals)
90     minbv = BitVector(intVal = min(intvals), size = P)
91     #print('min bit vector=',minbv)
92     bvruns = minbv.runs()
93     #print('bit vector runs=',bvruns)
94     if (len(bvruns) > 2): return P + 1
95     elif (len(bvruns) == 1 and bvruns[0][0] == '1'): return P
96     elif (len(bvruns) == 1 and bvruns[0][0] == '0'): return 0
97     else: return len(bvruns[1])
98
99 def get_lbp_hist(img_path, radius=1, num_neighbours=8):
100     P=num_neighbours
101     R=radius
102     img = cv2.imread(img_path,0)
103     img = cv2.resize(img,(256,256))
104     #img = cv2.resize(img,(40,40))
105     #print(img.shape)
106     hist =[0]*(P+2)
107     for x_idx in range(R, img.shape[1]-R):
108         for y_idx in range(R, img.shape[0]-R):
109             window = img[y_idx-1:y_idx+2, x_idx-1:x_idx+2]
110             pvalue = lbp(window)
111             hist[pvalue] += 1

```

```

112 hist = np.array(hist).astype(float) / np.sum(hist)
113 return hist
114
115 #get 1 image from each class and plot the histogram
116
117 im1_path = '/content/drive/My Drive/imagesDatabaseHW7/training/beach/1.jpg'
118 im1 = cv2.imread(im1_path)
119 im2 = cv2.cvtColor(im1, cv2.COLOR_BGR2RGB)
120 plt.figure(figsize=(10,10))
121 plt.imshow(im2)
122
123 hist = get_lbp_hist(im1_path)
124 plt.figure(figsize=(10,10))
125 fig=plt.bar(range(10), hist, edgecolor='black')
126 plt.xlabel('bins')
127 plt.ylabel('normalized frequency')
128 cv2.imwrite("beach.jpg", im1)
129 #cv2.imwrite("beach_hist.jpg", fig)
130 plt.savefig("beach_hist.jpg")
131
132
133 im1_path = '/content/drive/My Drive/imagesDatabaseHW7/training/building/01.jpg'
134
135 im1 = cv2.imread(im1_path)
136 im2 = cv2.cvtColor(im1, cv2.COLOR_BGR2RGB)
137 plt.figure(figsize=(10,10))
138 plt.imshow(im2)
139
140
141 hist = get_lbp_hist(im1_path)
142 plt.figure(figsize=(10,10))
143 fig=plt.bar(range(10), hist, edgecolor='black')
144 plt.xlabel('bins')
145 plt.ylabel('normalized frequency')
146 cv2.imwrite("building.jpg", im1)
147 plt.savefig("building_hist.jpg")
148
149 im1_path = '/content/drive/My Drive/imagesDatabaseHW7/training/mountain/01.jpg'
150
151 im1 = cv2.imread(im1_path)
152 im2 = cv2.cvtColor(im1, cv2.COLOR_BGR2RGB)
153 plt.figure(figsize=(10,10))
154 plt.imshow(im2)
155
156 hist = get_lbp_hist(im1_path)
157 plt.figure(figsize=(10,10))
158 fig=plt.bar(range(10), hist, edgecolor='black')
159 plt.xlabel('bins')
160 plt.ylabel('normalized frequency')
161 cv2.imwrite("mountain.jpg", im1)
162 plt.savefig("mountain_hist.jpg")
163
164
165 im1_path = '/content/drive/My Drive/imagesDatabaseHW7/training/tree/01.jpg'
166 im1 = cv2.imread(im1_path)
167 im2 = cv2.cvtColor(im1, cv2.COLOR_BGR2RGB)
168 plt.figure(figsize=(10,10))

```

```

168 plt.imshow(im2)
169
170 hist = get_lbp_hist(im1_path)
171 plt.figure(figsize=(10,10))
172 fig=plt.bar(range(10),hist,edgecolor='black')
173 plt.xlabel('bins')
174 plt.ylabel('normalized frequency')
175 cv2.imwrite("tree.jpg",im1)
176 plt.savefig("tree_hist.jpg")
177
178 im1_path = '/content/drive/My Drive/imagesDatabaseHW7/training/car/01.jpg'
179 im1 = cv2.imread(im1_path)
180 im2 = cv2.cvtColor(im1, cv2.COLOR_BGR2RGB)
181 plt.figure(figsize=(10,10))
182 plt.imshow(im2)
183
184 hist = get_lbp_hist(im1_path)
185 plt.figure(figsize=(10,10))
186 fig=plt.bar(range(10),hist,edgecolor='black')
187 plt.xlabel('bins')
188 plt.ylabel('normalized frequency')
189 cv2.imwrite("car.jpg",im1)
190 plt.savefig("car_hist.jpg")
191
192
193
194 # do this for every image in the folder
195 def train():
196     ##### loading images from folder code taken from :https://www.codegrepper.com
197     /code-examples/objectivec/how+to+load+images+from+folder+in+python
198     #training_dir='/content/drive/My Drive/imagesDatabaseHW7/training/'+
199     classname
200     training_dir='/content/drive/My Drive/imagesDatabaseHW7/training/'
201     classnames = os.listdir(training_dir)
202     i=0
203     for classname in classnames:
204         #print(classnames)
205         images = []
206         hist_list=[]
207         count=0
208         image_dir = training_dir+classname
209         for filename in os.listdir(image_dir):
210             hist = get_lbp_hist(os.path.join(image_dir,filename))
211             count+=1
212             print("histogram of " + classname + " number: "+str(count))
213             hist_list.append((hist,i))
214         with open(classname+'_features.pickle','wb') as fp:
215             pickle.dump(hist_list,fp)
216
217         with open(classname+'_features.pickle','rb') as fp:
218             pickle.load(fp)
219         i+=1
220     #return hist_list
221
222 train()
223
224 '''
225 class labels:

```

```

224     tree=0
225     beach=1
226     mountain=2
227     building=3
228     car=4
229     , , ,
230 with open('tree_features.pickle', 'rb') as fp:
231     tree_features = pickle.load(fp)
232 with open('beach_features.pickle', 'rb') as fp:
233     beach_features = pickle.load(fp)
234 with open('mountain_features.pickle', 'rb') as fp:
235     mountain_features = pickle.load(fp)
236 with open('building_features.pickle', 'rb') as fp:
237     building_features = pickle.load(fp)
238 with open('car_features.pickle', 'rb') as fp:
239     car_features = pickle.load(fp)
240 #print(beach_features[0][-1])
241
242 train_features = tree_features + beach_features + mountain_features +
    building_features + car_features
243
244 def test():
245     ##### code reference - Naveen Madapana (best solution1 2018) #####
246     testing_dir='/content/drive/My Drive/imagesDatabaseHW7/testing'
247     test_img_paths = glob.glob(os.path.join(testing_dir, '*.jpg'))
248     out_features = {os.path.basename(test_img_path): None for test_img_path in
        test_img_paths}
249     #print(out_features)
250     for test_img_path in test_img_paths:
251         out_feat_vec = get_lbp_hist(test_img_path)
252         out_features[os.path.basename(test_img_path)] = out_feat_vec
253     with open('test_features.pickle', 'wb') as fp:
254         pickle.dump(out_features, fp)
255
256     with open('test_features.pickle', 'rb') as fp:
257         out_features = pickle.load(fp)
258         class_dict = {"tree":0, "beach":1, "mountain":2, "building":3, "car":4}
259         test_features = []
260         for fname, test_inst in out_features.items():
261             cname = os.path.splitext(fname)[0].split('_')[0]
262             c_idx = class_dict[cname]
263             test_features.append((test_inst, c_idx))
264
265     return test_features
266
267 test_features=test()
268
269 def most_frequent(List):
270     #https://www.geeksforgeeks.org/python-find-most-frequent-element-in-a-list/
271     return max(set(List), key = List.count)
272
273 def get_confusion_matrix(pred, true, count):
274     conf_mat = np.zeros((5,5))
275     correct=0
276     for i in range(len(true)):
277         if pred[i]==true[i]:
278             correct+=1

```

```

279     conf_mat[true[i],pred[i]]+=1
280     accuracy = correct/len(pred)
281     print("Accuracy: ",accuracy)
282     print()
283     print("confusion matrix:\n",conf_mat)
284
285
286 def knn(train_input , test_input , metric='Euclidean' ,K=5):
287     train_input = np.array(train_input)
288     test_input = np.array(test_input)
289     train_data = train_input[:,0]
290     train_label = train_input[:,1]
291     test_data = test_input[:,0]
292     test_label = test_input[:,1]
293     pred_labels=[]
294     for i in range(len(test_input)):
295         dist =[]
296         for j in range(len(train_input)):
297             if metric == 'Euclidean':
298                 distance = np.linalg.norm(train_data[j]-test_data[i])
299             elif metric == 'Manhattan':
300                 distance = np.sum(np.abs(train_data[j]-test_data[i]))
301                 #distance =0
302             dist.append((distance , train_label[j]))
303             #print(dist)
304         dist.sort(key=lambda x:x[0])
305         #print(dist)
306         labels=[ x[1] for x in dist]
307         topK=labels[:K]
308         pred_labels.append(most_frequent(topK))
309
310     get_confusion_matrix(pred_labels , test_label , len(test_label))
311     return pred_labels
312
313 labels =knn(train_features , test_features , metric='Manhattan' ,K=7)

```